

```

In [3]: # =====
# Task 3: Linear Regression (Auto-detect version)
# =====

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# =====
# 1. Load Dataset
# =====
# Replace with your dataset filename
df = pd.read_csv("Housing.csv")

print("Dataset Shape:", df.shape)
print("First 5 rows:\n", df.head())
print("\nAvailable columns:\n", df.columns)

# Drop missing values
df = df.dropna()

# =====
# 2. Auto Feature & Target Selection
# =====
# Keep only numeric columns
numeric_cols = df.select_dtypes(include=[np.number]).columns.tolist()
print("\nNumeric columns found:", numeric_cols)

# Assume last numeric column = target (Price)
target_col = numeric_cols[-1]
feature_cols = numeric_cols[:-1]

print("Selected Features:", feature_cols)
print("Selected Target:", target_col)

X = df[feature_cols]
y = df[target_col]

# =====
# 3. Train-Test Split
# =====
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# =====
# 4. Train Model
# =====
model = LinearRegression()
model.fit(X_train, y_train)

# =====
# 5. Predictions
# =====
y_pred = model.predict(X_test)

# =====
# 6. Evaluation
# =====
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

print("\nModel Performance:")
print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
print("R² Score:", r2)

# =====
# 7. Coefficients
# =====
print("\nIntercept:", model.intercept_)
coeff_df = pd.DataFrame({
    "Feature": feature_cols,
    "Coefficient": model.coef_
})

```

```

})
print("\nFeature Coefficients:\n", coeff_df)

# =====
# 8. Plot Regression (if simple regression)
# =====
if len(feature_cols) == 1:
    plt.scatter(X_test, y_test, color="blue", label="Actual")
    plt.plot(X_test, y_pred, color="red", linewidth=2, label="Predicted")
    plt.xlabel(feature_cols[0])
    plt.ylabel(target_col)
    plt.title("Simple Linear Regression")
    plt.legend()
    plt.show()

# =====
# 9. Residual Plot (for checking assumptions)
# =====
residuals = y_test - y_pred
sns.scatterplot(x=y_pred, y=residuals)
plt.axhline(y=0, color="red", linestyle="--")
plt.xlabel("Predicted Values")
plt.ylabel("Residuals")
plt.title("Residual Plot")
plt.show()

```

Dataset Shape: (545, 13)

First 5 rows:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	\
0	13300000	7420	4	2	3	yes	no	no	
1	12250000	8960	4	4	4	yes	no	no	
2	12250000	9960	3	2	2	yes	no	yes	
3	12215000	7500	4	2	2	yes	no	yes	
4	11410000	7420	4	1	2	yes	yes	yes	

	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
0	no	yes	2	yes	furnished
1	no	yes	3	no	furnished
2	no	no	2	yes	semi-furnished
3	no	yes	3	yes	furnished
4	no	yes	2	no	furnished

Available columns:

```

Index(['price', 'area', 'bedrooms', 'bathrooms', 'stories', 'mainroad',
      'guestroom', 'basement', 'hotwaterheating', 'airconditioning',
      'parking', 'prefarea', 'furnishingstatus'],
      dtype='object')

```

Numeric columns found: ['price', 'area', 'bedrooms', 'bathrooms', 'stories', 'parking']

Selected Features: ['price', 'area', 'bedrooms', 'bathrooms', 'stories']

Selected Target: parking

Model Performance:

MAE: 0.6334119253074165

MSE: 0.6078468389300876


RMSE: 0.7796453289349508

R<sup>2</sup> Score: 0.2272813724236712

Intercept: -0.3348998368179409

Feature Coefficients:

	Feature	Coefficient
0	price	1.500301e-07
1	area	6.652296e-05
2	bedrooms	4.320597e-02
3	bathrooms	2.571284e-02
4	stories	-1.058846e-01



In [ ]: